

Model Comparison: trinity vs qwen3.6

Date: 2026-05-04

This document compares trinity and qwen3.6 using 2 independent evaluation sources. Each source may evaluate different model variants, so the sections should be read as complementary views rather than direct cross-references.

Section	Models Found	Evaluation Method
Part 1: Arena	qwen3.6-max-preview, qwen3.6-plus, trinity-large-preview, trinity-large-thinking	Human Preference Ratings
Part 2: Artificial Analysis	Trinity Large Thinking, Qwen3.6 27B (Reasoning), Qwen3.6 35B A3B (Reasoning), Qwen3.6 27B (Non-reasoning), Qwen3.6 Plus, Qwen3.6 35B A3B (Non-reasoning), Qwen3.6 Max Preview	Automated Benchmarks

Overall Conclusions

- Overall positioning:** Qwen3.6 operates at the Near-frontier to Upper-mid tier, while Trinity sits squarely in Mid-tier territory. In Arena, Qwen3.6 Max Preview ranks 31st of 356 models (rating 1456.7), placing it alongside Claude Sonnet 4.5, GPT-5.1, and Gemini 2.5 Pro. Trinity Large Thinking ranks 133rd (rating 1382.2), near GPT-4.1 Mini and Mistral Medium. In AA benchmarks, the gap is equally stark: Qwen3.6 Max Preview scores 51 on the Intelligence Index (rank 16 of 362), while Trinity Large Thinking scores 31 (rank 129). Both sources show clear separation between the two families.
- Lineup depth:** Qwen3.6 offers a broader lineup: 2 API-tier models (Max Preview, Plus) and 4 open-weight variants in AA (27B and 35B in both reasoning and non-reasoning modes). Trinity has just 2 models total: trinity-large-thinking and trinity-large-preview, with only the thinking variant benchmarked in AA. Qwen3.6's lineup spans from the Near-frontier (Intelligence 51) down to Mid-tier (Intelligence 31), covering a wider range of deployment needs.
- Value proposition:** Trinity's strongest case is operational efficiency. On AA benchmarks, Trinity Large Thinking is 3.7x faster than Qwen3.6 Max Preview (141 vs 38 t/s), has 4.2x lower latency (0.49s vs 2.04s TTFT), and is 7.4x cheaper (\$0.40 vs \$2.92/1M blended tokens). Even against the more affordable Qwen3.6 Plus (\$1.12/1M), Trinity is still 2.8x cheaper. However, this cost advantage comes with a significant quality gap. For users who need maximum quality, Qwen3.6 is the clear choice; for latency-sensitive, budget-constrained applications, Trinity offers a compelling speed/cost trade-off.
- Quality profile differences:** In Arena's head-to-head data, Qwen3.6 has a STEM-leaning advantage profile, with its largest margins in math (+106.9 avg) and industrial mathematics (+103.8 avg). Its smallest advantage is in coding (+65.6 avg), where Trinity's relative strength shows through — Trinity Large Preview actually scores 1442.2 in coding, competitive with many higher-ranked models. Creative writing shows a moderate gap (+66.1 avg), suggesting Trinity's conversational and creative capabilities are somewhat closer to Qwen3.6 than its STEM capabilities.
- Tier-matched comparison:** At the same intelligence level, Trinity holds its own. In AA, Trinity Large Thinking (Intelligence 31, Coding 27) matches Qwen3.6 35B A3B Non-reasoning (Intelligence 31, Coding 17) while delivering higher coding scores, faster inference (141 vs 189 t/s), lower latency (0.49s vs 1.36s), and a lower price (\$0.40 vs \$0.84/1M). This suggests that within its tier, Trinity is a strong competitor — the quality gap only appears when comparing against Qwen3.6's top-tier API offerings.
- Evaluation coverage:** Trinity has limited evaluation coverage: only 1 model in AA and 2 in Arena. Qwen3.6 has 6 models in AA and 2 in Arena. Neither source reports context window or parameter

counts for these models, limiting architecture-level comparisons. The AA comparison necessarily compares Qwen3.6’s top model (a proprietary API offering) against Trinity’s single evaluated model, which may not be a fair apples-to-apples comparison.

7. Summary:

Factor	Trinity	Qwen3.6
Top-tier quality (Arena)	1382 (rank 133, Mid-tier)	1457 (rank 31, Upper-mid)
Top-tier quality (AA)	Intelligence 31 (rank 129)	Intelligence 51 (rank 16)
Same-tier quality	Intelligence 31, Coding 27	Intelligence 31, Coding 17 (35B A3B)
Speed	141 t/s	38–199 t/s (varies by model)
Latency (TTFT)	0.49s	1.36–2.04s
Price	\$0.40/1M blend	\$0.56–\$2.92/1M blend
Coding (Arena)	1430–1442	1498–1506
Math (Arena)	1356–1396	1478–1488
Model variety	2 models	6 models (AA), 2 models (Arena)
Open weights	Yes (Arcee)	Yes (27B, 35B variants)

8. **Bottom line:** Qwen3.6 is the higher-quality family by a wide margin, operating 1–2 tiers above Trinity across both human preference and automated benchmarks. Choose Qwen3.6 when quality is the priority — its Max Preview and Plus models compete with the best in the industry. Choose Trinity when you need a fast, affordable, open-weight model with respectable coding ability and can accept mid-tier quality: at \$0.40/1M tokens and sub-500ms TTFT, it’s well-suited for latency-sensitive pipelines, high-volume inference, or as a cost-effective component in multi-model architectures.

Part 1: Arena

Source: Arena

About Arena


Arena (formerly LMSYS Chatbot Arena) ranks models using **head-to-head human preference votes**. Real users submit prompts to two anonymous models side-by-side, then choose which response they prefer. These pairwise outcomes are aggregated into Bradley-Terry ratings (similar to Elo in chess) – higher is better, with scores typically ranging from ~900 to ~1550.

The `text_style_control` variant adjusts for verbosity bias, so models don’t get rewarded simply for producing longer responses.

Ratings are broken down into **27 topic-based categories** (coding, math, creative writing, legal, etc.) based on conversation content, giving a detailed profile of where each model excels relative to the full field. Because these are human judgments rather than automated test suites, they reflect what real users find helpful – but they also carry the biases of the voter population (skewed toward AI-savvy early adopters) and the preference-vs-correctness gap (a confident but wrong answer can still win votes).

Models evaluated: qwen3.6-max-preview, qwen3.6-plus, trinity-large-preview, trinity-large-thinking

Global Arena Rankings (356 models total)



Rank	Model	Rating	Votes
26	gemini-3-flash (thinking-minimal)	1462.9	40,022
27	deepseek-v4-pro-thinking	1461.7	3,820
28	kimi-k2.6	1460.0	6,168
29	grok-4.1	1459.9	57,888
30	dola-seed-2.0-pro	1459.4	25,273
31	qwen3.6-max-preview	1456.7	3,943
32	glm-5	1456.3	19,980
33	gpt-5.4-mini-high	1456.1	13,541
34	grok-4.3	1455.7	2,505
35	gpt-5.1-high	1454.6	40,879
36	claude-sonnet-4-5-20250929- thinking-32k	1453.6	65,936
37	claude-sonnet-4-5-20250929	1452.6	63,899
38	gemma-4-31b	1451.0	5,810
39	ernie-5.0-0110	1450.7	27,672
40	kimi-k2.5-thinking	1449.5	26,123
41	ernie-5.0-preview-1203	1449.2	9,770
42	mimo-v2-pro	1449.1	14,085
43	claude-opus-4-1-20250805- thinking-16k	1448.7	49,853
44	gpt-5.3-chat-latest	1448.4	21,109
45	gemini-2.5-pro	1447.7	113,545
46	qwen3.6-plus	1447.3	7,538
47	qwen3.5-397b-a17b	1447.1	21,382
48	claude-opus-4-1-20250805	1446.7	77,437
49	gpt-4.5-preview-2025-02-27	1444.2	14,547
50	chatgpt-4o-latest-20250326	1442.7	82,545
51	glm-4.7	1442.5	12,143
[76 models not shown]			
128	claude-3-7-sonnet-20250219- thinking-32k	1386.7	38,834

Rank	Model	Rating	Votes
129	mistral-medium-2505	1386.4	33,248
130	minimax-m2.1-preview	1385.1	17,157
131	qwen3-30b-a3b-instruct-2507	1383.2	23,760
132	hunyuan-turbos-20250416	1382.4	10,720
133	trinity-large-thinking	1382.2	10,875
134	gpt-4.1-mini-2025-04-14	1382.2	39,360
135	gemini-2.5-flash-lite-preview-09-2025-no-thinking	1379.6	47,285
136	glm-4.6v	1377.6	2,807
137	trinity-large-preview	1374.6	19,586
138	qwen3-235b-a22b	1374.5	26,280
139	gemini-2.5-flash-lite-preview-06-17-thinking	1374.4	32,944
140	qwen2.5-max	1374.1	32,623
141	glm-4.5-air	1372.8	31,126
142	claude-3-5-sonnet-20241022	1371.9	88,358

Subset Rankings (4 models)

Rank	Model	Votes	Overall
1	qwen3.6-max-preview	3,943	1456.7
2	qwen3.6-plus	7,538	1447.3
3	trinity-large-thinking	10,875	1382.2
4	trinity-large-preview	19,586	1374.6

Part 1: General capabilities

Model	overall	coding	math	creative	instruct	hard
qwen3.6-max-preview	1456.7	1497.7	1487.8	1432.0	1448.6	1477.9
qwen3.6-plus	1447.3	1505.8	1478.1	1397.1	1434.3	1468.2
trinity-large-thinking	1382.2	1430.1	1396.0	1347.5	1370.2	1398.4
trinity-large-preview	1374.6	1442.2	1356.1	1349.4	1368.1	1400.4

Part 2: Conversational and industry categories

Model	multi_turn	ind_math
qwen3.6-max-preview	1469.3	1510.6
qwen3.6-plus	1438.8	1453.7
trinity-large-thinking	1357.9	1386.2
trinity-large-preview	1374.2	1370.4

Win/Loss Summary

vs Model	qwen3.6-max-preview Wins	Opponent Wins	Overall Gap
trinity-large-preview	8	0	+82.2
trinity-large-thinking	8	0	+74.5
trinity-large-preview	8	0	+72.8
trinity-large-thinking	8	0	+65.1

Head-to-Head: qwen3.6-max-preview vs trinity-large-thinking

qwen3.6-max-preview wins 8 of 8 categories.

Category	qwen3.6-max-preview	trinity-large-thinking	Delta	Winner
ind_math	1510.6	1386.2	+124.4	qwen3.6-max-preview
multi_turn	1469.3	1357.9	+111.4	qwen3.6-max-preview
math	1487.8	1396.0	+91.8	qwen3.6-max-preview
creative	1432.0	1347.5	+84.5	qwen3.6-max-preview
hard	1477.9	1398.4	+79.6	qwen3.6-max-preview
instruct	1448.6	1370.2	+78.4	qwen3.6-max-preview
overall	1456.7	1382.2	+74.5	qwen3.6-max-preview
coding	1497.7	1430.1	+67.6	qwen3.6-max-preview

Head-to-Head: qwen3.6-max-preview vs trinity-large-preview

qwen3.6-max-preview wins 8 of 8 categories.

Category	qwen3.6-max-preview	trinity-large-preview	Delta	Winner
ind_math	1510.6	1370.4	+140.2	qwen3.6-max-preview
math	1487.8	1356.1	+131.7	qwen3.6-max-preview
multi_turn	1469.3	1374.2	+95.1	qwen3.6-max-preview
creative	1432.0	1349.4	+82.6	qwen3.6-max-preview
overall	1456.7	1374.6	+82.2	qwen3.6-max-preview
instruct	1448.6	1368.1	+80.5	qwen3.6-max-preview
hard	1477.9	1400.4	+77.5	qwen3.6-max-preview

Category	qwen3.6-max-preview	trinity-large-preview	Delta	Winner
coding	1497.7	1442.2	+55.5	qwen3.6-max-preview

Head-to-Head: qwen3.6-plus vs trinity-large-thinking

qwen3.6-plus wins 8 of 8 categories.

Category	qwen3.6-plus	trinity-large-thinking	Delta	Winner
math	1478.1	1396.0	+82.2	qwen3.6-plus
multi_turn	1438.8	1357.9	+80.9	qwen3.6-plus
coding	1505.8	1430.1	+75.7	qwen3.6-plus
hard	1468.2	1398.4	+69.8	qwen3.6-plus
ind_math	1453.7	1386.2	+67.4	qwen3.6-plus
overall	1447.3	1382.2	+65.1	qwen3.6-plus
instruct	1434.3	1370.2	+64.1	qwen3.6-plus
creative	1397.1	1347.5	+49.6	qwen3.6-plus

Head-to-Head: qwen3.6-plus vs trinity-large-preview

qwen3.6-plus wins 8 of 8 categories.

Category	qwen3.6-plus	trinity-large-preview	Delta	Winner
math	1478.1	1356.1	+122.1	qwen3.6-plus
ind_math	1453.7	1370.4	+83.2	qwen3.6-plus
overall	1447.3	1374.6	+72.8	qwen3.6-plus
hard	1468.2	1400.4	+67.8	qwen3.6-plus
instruct	1434.3	1368.1	+66.2	qwen3.6-plus
multi_turn	1438.8	1374.2	+64.6	qwen3.6-plus
coding	1505.8	1442.2	+63.5	qwen3.6-plus
creative	1397.1	1349.4	+47.7	qwen3.6-plus

Arena Key Findings

1. **Overall positioning:** Qwen3.6 Max Preview (rank 31 of 356, rating 1456.7) sits in the Upper-mid tier alongside models like Claude Sonnet 4.5, GPT-5.1, and Gemini 2.5 Pro — strong company that indicates genuine frontier-adjacent quality. Trinity Large Thinking (rank 133, rating 1382.2) sits in the Mid-tier, near GPT-4.1 Mini and Claude 3.7 Sonnet (non-thinking). The 74.5-point gap between these two represents clear separation — their confidence intervals do not overlap, and the gap exceeds a full standard deviation of the Arena population.
2. **Head-to-head summary:** Qwen3.6 sweeps all 8 evaluated categories in every one of the 4 cross-family matchups, leaving no category where Trinity holds an advantage. This is unusually comprehensive dominance and reflects the tier gap between the two families.
3. **Qwen3.6’s largest advantages are in mathematical reasoning**, where the gap is largest: math (+106.9 avg, peaking at +131.7 vs trinity-large-preview) and industrial mathematics (+103.8 avg, peaking at +140.2). These are among the largest category deltas in the comparison, suggesting Qwen3.6’s architecture and training are particularly optimized for quantitative tasks. Multi-turn conversation also shows a large gap (+88.0 avg), indicating Qwen3.6 maintains quality better over extended interactions.

4. **Trinity’s relative strength is coding**, where the gap is smallest (+65.6 avg). Trinity Large Preview scores 1442.2 in coding — higher than its thinking variant (1430.1) and competitive with models ranked 20+ positions above it overall. This suggests coding is Trinity’s standout capability relative to its overall tier placement, and users primarily focused on code generation may find the quality gap more tolerable.
 5. **Profile difference:** Qwen3.6’s advantage is STEM-leaning (avg delta +92.1 in STEM categories vs +69.2 in humanities/creative). While Qwen3.6 dominates everywhere, the gap narrows in creative writing (+66.1 avg) and instruction following (+72.3 avg), areas where Trinity is relatively more competitive. Users whose workloads skew toward writing and conversation may see a somewhat smaller practical gap than the overall rankings suggest.
-

Part 2: Artificial Analysis

Source: Artificial Analysis

About Artificial Analysis

Artificial Analysis evaluates models using **automated benchmark suites** – standardized tests with known correct answers, run programmatically. Their Intelligence Index (v4.0) aggregates scores from 10 evaluations:


- **GPQA Diamond** – graduate-level science questions
- **Humanity’s Last Exam (HLE)** – extremely difficult cross-domain questions
- **SciCode** – scientific coding problems
- **Terminal-Bench Hard** – complex terminal/CLI tasks
- **IFBench** – instruction following
- **AA-LCR** – long-context retrieval
- **AA-Omniscience** – broad knowledge assessment
- **GDPval-AA** – GDP prediction (quantitative reasoning)
- **tau2-Bench Telecom** – domain-specific agent tasks
- **CritPt** – critical thinking

Unlike Arena’s human preference votes, these benchmarks have **objectively correct answers**. This makes AA scores more precise for measurable capabilities (coding, math, factual recall) but less reflective of subjective qualities like writing style, helpfulness, or conversational fluency.

AA also independently measures **speed** (output tokens/sec), **latency** (time to first token), and **pricing** across API providers, providing a practical deployment perspective.

Models evaluated: Trinity Large Thinking, Qwen3.6 27B (Reasoning), Qwen3.6 35B A3B (Reasoning), Qwen3.6 27B (Non-reasoning), Qwen3.6 Plus, Qwen3.6 35B A3B (Non-reasoning), Qwen3.6 Max Preview

Global AA Rankings (362 models total)



Rank	Model	AA Intelligence
11	Claude Opus 4.6 (Adaptive Reasoning, Max Effort)	53
12	Claude Sonnet 4.6 (Adaptive Reasoning, Max Effort)	51
13	Claude Opus 4.7 (Non-reasoning, High Effort)	51
14	DeepSeek V4 Pro (Reasoning, Max Effort)	51
15	GLM-5.1 (Reasoning)	51
16	Qwen3.6 Max Preview	51
17	GPT-5.2 (xhigh)	51
18	GPT-5.5 (low)	50
19	Qwen3.6 Plus	50
20	DeepSeek V4 Pro (Reasoning, High Effort)	49
21	MiniMax-M2.7	49
22	MiMo-V2-Pro	49
23	GLM-5 (Reasoning)	49
24	GPT-5.2 Codex (xhigh)	49
25	Claude Opus 4.5 (Reasoning)	49
26	Grok 4.20 0309 v2 (Reasoning)	49
27	GPT-5.4 mini (xhigh)	48
28	Gemini 3 Pro Preview (high)	48
29	Grok 4.20 0309 (Reasoning)	48
30	GPT-5.4 (low)	47
31	GPT-5.1 (high)	47
32	Gemini 3 Flash Preview (Reasoning)	46
33	DeepSeek V4 Flash (Reasoning, Max Effort)	46
34	GPT-5.2 (medium)	46
35	Claude Opus 4.6 (Non-reasoning, High Effort)	46
36	Kimi K2.5 (Reasoning)	46
37	Qwen3.6 27B (Reasoning)	45
38	Qwen3.5 397B A17B (Reasoning)	45
39	GPT-5.4 nano (xhigh)	44
40	Claude Sonnet 4.6 (Non-reasoning, High Effort)	44
41	DeepSeek V4 Flash (Reasoning, High Effort)	44
42	MiMo-V2-Omni-0327	44

Rank	Model	AA Intelligence
43	GPT-5 (high)	44
44	GPT-5 Codex (high)	44
45	Kimi K2.6 (Non-reasoning)	43
46	MiMo-V2-Omni	43
47	KAT Coder Pro V2	43
48	GLM-5.1 (Non-reasoning)	43
49	Qwen3.6 35B A3B (Reasoning)	43
50	GPT-5.1 Codex (high)	43
51	Claude 4.5 Sonnet (Reasoning)	43
52	Claude Opus 4.5 (Non-reasoning)	43
53	Claude Sonnet 4.6 (Non-reasoning, Low Effort)	42
54	GPT-5 (medium)	42
	<i>[30 models not shown]</i>	
85	GPT-5 mini (medium)	38
86	Claude 4 Sonnet (Reasoning)	38
87	GPT-5.4 mini (medium)	37
88	Claude 4.5 Haiku (Reasoning)	37
89	Kimi K2.5 (Non-reasoning)	37
90	Qwen3.6 27B (Non-reasoning)	37
91	Claude 4.5 Sonnet (Non-reasoning)	37
92	Step 3.5 Flash	37
93	Qwen3.5 35B A3B (Reasoning)	37
94	Qwen3.5 27B (Non-reasoning)	37
95	DeepSeek V4 Flash (Non-reasoning)	36
	<i>[28 models not shown]</i>	
124	GLM-4.6 (Reasoning)	32
125	Qwen3 Max Thinking (Preview)	32
126	Gemma 4 26B A4B (Reasoning)	31
127	Claude 4.5 Haiku (Non-reasoning)	31
128	Nova 2.0 Pro Preview (low)	31
129	Trinity Large Thinking	31
130	Qwen3.6 35B A3B (Non-reasoning)	31
131	Qwen3 Max	31
132	MiMo-V2-Flash (Non-reasoning)	30
133	Qwen3.5 35B A3B (Non-reasoning)	30
134	o1	30
135	Claude 3.7 Sonnet (Non-reasoning)	30

Reasoning Models

Model	Params (total/active)	AA Intelli- gence	Coding	Speed (t/s)	TTFT (s)	Price (\$/1M blend)	Context
Qwen3.6 27B (Reasoning)	proprietary	45	36	65.0	1.52	\$1.35	–
Qwen3.6 35B A3B (Reasoning)	proprietary	43	35	198.8	1.45	\$0.56	–

Model	Params (total/active)	AA Intelligence	Coding	Speed (t/s)	TTFT (s)	Price (\$/1M blend)	Context
Qwen3.6 27B (Non-reasoning)	proprietary	37	26	67.5	1.58	\$1.35	–
Trinity Large Thinking	proprietary	31	27	140.7	0.49	\$0.40	–
Qwen3.6 35B A3B (Non-reasoning)	proprietary	31	17	188.7	1.36	\$0.84	–

Non-Reasoning Models

Model	Params (total/active)	AA Intelligence	Coding	Speed (t/s)	TTFT (s)	Price (\$/1M blend)	Context
Qwen3.6 Max Preview	proprietary	51	44	37.6	2.04	\$2.92	–
Qwen3.6 Plus	proprietary	50	42	52.8	1.75	\$1.12	–

All Models

Model	AA Intelligence	Coding	Speed (t/s)	TTFT (s)	Price (\$/1M blend)	Context
Qwen3.6 Max Preview	51	44	37.6	2.04	\$2.92	–
Qwen3.6 Plus	50	42	52.8	1.75	\$1.12	–
Qwen3.6 27B (Reasoning)	45	36	65.0	1.52	\$1.35	–
Qwen3.6 35B A3B (Reasoning)	43	35	198.8	1.45	\$0.56	–
Qwen3.6 27B (Non-reasoning)	37	26	67.5	1.58	\$1.35	–
Trinity Large Thinking	31	27	140.7	0.49	\$0.40	–
Qwen3.6 35B A3B (Non-reasoning)	31	17	188.7	1.36	\$0.84	–

Artificial Analysis Key Findings

1. **Arcee AI positioning:** Trinity Large Thinking scores 31 on the Intelligence Index (~rank 129 of 362), placing it in the Mid-tier alongside models like Claude 4.5 Haiku (Non-reasoning), Gemma 4 26B A4B, and older Qwen3 Max. Only 1 Trinity model is evaluated in AA, limiting lineup depth comparisons.
2. **Alibaba positioning:** Qwen3.6 offers a strong lineup with 6 evaluated models spanning from Intelligence 51 (Max Preview, rank 16) down to Intelligence 31 (35B A3B Non-reasoning, rank 130). The

top two models — Max Preview and Plus — are Near-frontier, sitting alongside Claude Opus 4.7, DeepSeek V4 Pro, and GLM-5.1. Even the mid-range Qwen3.6 27B (Reasoning) at Intelligence 45 outperforms many competitors’ flagship offerings.

3. **Intelligence gap:** Qwen3.6 Max Preview (51) vs Trinity Large Thinking (31) is a 20-point gap — clear separation representing over 1.5 standard deviations of the AA population. To put this in perspective, the entire range from the top model (GPT-5.5 xhigh at 60) to Trinity spans 29 points, meaning Qwen3.6 Max Preview captures roughly 69% of the gap to the absolute best. However, comparing Qwen3.6’s top API model against Trinity’s single offering may overstate the practical gap — Qwen3.6 35B A3B (Non-reasoning) matches Trinity at Intelligence 31.
4. **Coding:** Qwen3.6 Max Preview leads with Coding Index 44 vs Trinity’s 27 — a 17-point gap representing clear separation. However, Trinity’s Coding Index of 27 exceeds Qwen3.6 35B A3B Non-reasoning’s 17, suggesting that at the same intelligence tier, Trinity has a coding advantage. This aligns with Trinity’s relatively strong coding performance seen in Arena data.
5. **Speed advantage — Trinity:** Trinity Large Thinking is dramatically faster at 3.7x the throughput of Qwen3.6 Max Preview (141 vs 38 t/s). Even against the faster Qwen3.6 variants — the 35B A3B models run at 189–199 t/s — Trinity remains competitive. This speed makes Trinity well-suited for high-volume inference pipelines where throughput matters more than peak quality.
6. **Latency advantage — Trinity:** Trinity’s 0.49s TTFT is 4.2x faster than Qwen3.6 Max Preview (2.04s) and roughly 3x faster than even the lightest Qwen3.6 variants (~1.36–1.58s). Sub-500ms TTFT is a meaningful threshold for interactive applications, placing Trinity among the most responsive models in its quality band.
7. **Price advantage — Trinity:** At \$0.40/1M blended tokens, Trinity is 7.4x cheaper than Qwen3.6 Max Preview (\$2.92) and 2.8x cheaper than Qwen3.6 Plus (\$1.12). Only the Qwen3.6 35B A3B models (\$0.56–\$0.84) approach Trinity’s price point, and at that price the two families offer comparable intelligence (both at 31) with different coding strengths.

Definitions

Model Tiers

Tiers are based on absolute rank within each data source:

Tier	Rank Range
Frontier	1–10
Near-frontier	11–25
Upper-mid	26–75
Mid-tier	76–150
Long-tail	151+

Gap Significance

Arena — uses per-model confidence intervals (CI). If CIs overlap, the models are “statistically indistinguishable.” If CIs don’t overlap but the closest bounds are within 20 points, the difference is “small but statistically significant.” A gap of 20+ points between closest bounds indicates “clear separation.”

Artificial Analysis — uses gap relative to population standard deviation. A gap below 0.5 stdev is “not clearly distinguishable,” 0.5–1.0 stdev indicates a “moderate difference,” and above 1.0 stdev indicates “clear separation.”

Data provided by *Artificial Analysis* (<https://artificialanalysis.ai>) and *Arena* (<https://lmarena.ai>).